

Shayan Poigai

Fremont, CA | (510) 738-8155 | spoigai21@gmail.com | U.S. citizen
shayanpoigai.dev | linkedin.com/in/shayanpoigai | github.com/spoigai21

EDUCATION

Santa Clara University — B.S. in Computer Science

Expected December 2027

GPA: 3.9

Available Summer 2027 (Fall 2027 term remaining)

- **Relevant Coursework:** Operating Systems, Advanced C++, Advanced Algorithms, Data Structures & Algorithms, Theory of Automata, Discrete Mathematics

WORK EXPERIENCE

Software Engineering Intern

June 2026 – August 2026

HerbsPro

- Cut MySQL query volume 60% and product API p50 from 200 ms to 90 ms with a cache-aside Redis layer using event-driven invalidation across a 50,000+ SKU catalog.
- Built a Python MCP server (FastMCP) as the only access path for LLM agents into a legacy MySQL WMS: read tools for live stock, write tools emitting split-shipment pick tickets instead of raw table mutations.
- Built a Pandas **ETL pipeline** that normalized 150+ unstructured meeting notes into a fixed schema, extracting supplier terms into structured records.

Software Engineer

May 2026 – Present

Adorus

- Built the production platform from scratch, solo (530 commits): a Java 21 Spring Boot backend of **24 bounded modules**, a React + Vite storefront, and PostgreSQL on AWS behind CloudFront, owning every layer from schema to CDN.
- Shipped an **AI try-on feature** end to end: photo upload and live camera capture in React, asynchronous generation jobs, and a credit system, with a guarantee that no customer photo is ever persisted.
- Built a **92-module React + Vite storefront** across 41 indexable routes, with client-side routing, a guest cart on a 30-day httpOnly cookie, and static metadata for crawlers and social scrapers that never execute JavaScript.
- Shipped search by photo: a YOLO11-pose crop and CLIP embeddings moved the correct match from rank 3 to rank 1, with inference running in AWS Lambda.

AI/ML Research Assistant

Dec 2025 – June 2026

Santa Clara University

- Co-author and primary writer of *Public Perceptions of ChatGPT as a Therapist-Like Tool*, accepted at HI-BI-BI 2026 (ASONAM-collocated); proceedings forthcoming with Springer.
- Analyzed 5,000+ scraped Twitter, Reddit, and YouTube posts with an NLP and **machine learning** pipeline comparing BERTopic, LDA, and NMF to measure sentiment toward LLMs as a therapeutic tool.

Lab Teaching Assistant (C++)

Sept 2025 – June 2026

Santa Clara University

- Mentored 40+ students in C++ memory management, pointers, and object-oriented design, **debugging** logic and memory errors live with GDB.

PROJECTS

Internship Monitor | *Python, SQLite, GitHub Actions, pytest*

2026

- Co-built a Python job monitor tracking 106 companies across 17 job-board types; added the Simplify, Oracle, and SmartRecruiters parsers to a 308-test suite.
- Suppressed repeat alerts with a normalized employer-plus-title dedup key after finding ~2,900 of 14,808 listings were reposts issued under new job IDs.

Disease Tracker | *FastAPI, AWS Bedrock, DynamoDB*

Oct 2025

- Built the FastAPI backend for a 6-person team: disease-metric endpoints over DynamoDB and a county risk-scoring endpoint calling Claude on AWS Bedrock.

TECHNICAL SKILLS

Languages: Python, Java, JavaScript, SQL, C++

Frameworks & Tools: React.js (Vite), Node.js, Spring Boot, FastAPI, REST/RESTful APIs, Docker (containers), Linux, Git, GitHub Actions (CI/CD), nginx, Vercel, Render, Stripe

Cloud & Data: AWS (EC2, RDS, Bedrock, Cognito, S3, CloudFront, Lambda), PostgreSQL, MySQL, Redis, DynamoDB (NoSQL), ChromaDB

AI / ML: Machine Learning, NLP, computer vision, RAG, MCP, PyTorch, CLIP, vector embeddings, Ollama, OpenWebUI, NumPy, Pandas, Matplotlib, topic modeling (BERTopic, LDA, NMF)

Quantum: Qiskit, IBM Quantum